

ECE 661: Homework #4

Pruning and Fixed-point Quantization

NetID: wh162

1. True/False Questions

- a. Problem 1.1: True, weight pruning is a technique that discards unnecessary data in a model without affecting the accuracy, whereas weight optimization is how the weight is optimized for reducing the bit representation. So these two methods don't intervene with each other.
- b. Problem 1.2: False, even though the parameters are pruned to zero, the GPU still will compute the data when fed into it.
- c. Problem 1.3: False, though pruning will reduce the number of weights, the bit representation for a weight is still in 32/64 bit, which would take Huffman encoding more time to compute.
- d. Problem 1.4: False, pruning is to remove the weights that would less affect the accuracy of the model, it is not guaranteed that the value of the weight would be exact zero, so pruning is still necessary.
- e. Problem 1.5: False, as from the slides, though soft thresholding reveals the bias problem of L1, the issue was solved by SCAD and MCP.
- f. Problem 1.6: True, refer to lec13 page 21.
- g. Problem 1.7: True, refer to lec 14 page 16.

h. Problem 1.8: True, refer to lec 14 page 27-28

i. Problem 1.9: True, refer to lec15 page 10.

j. Problem 1.10: True, refer to Lec 15 page 28.


2. Lab1: Sparse optimization of linear model

a. Problem 2.1:

$$\begin{aligned} L &= (XW - Y)^2 \\ \frac{dL}{dW} &= \frac{d(X^2 W^2 - 2XWY + Y^2)}{dW} \\ &= 2X^2 W - 2XY \\ &= 2X(XW - Y) \end{aligned}$$

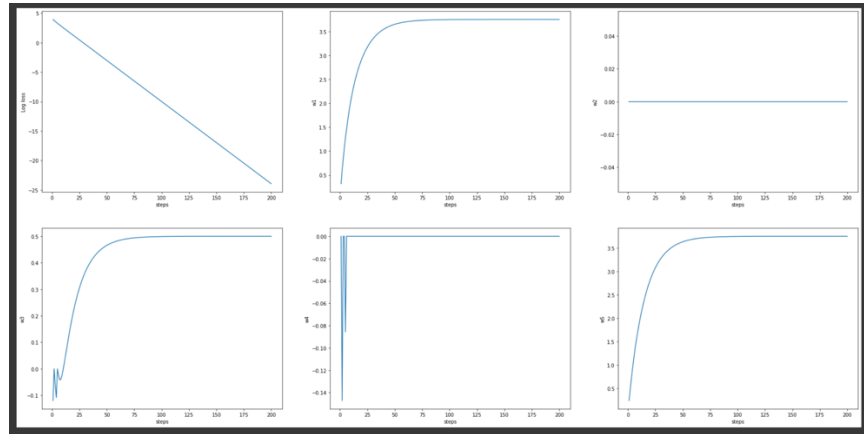
b. Problem 2.2:

From the plot below, we can see that W is converging to a value after several steps are executed. Yes, W is converging to an optimal solution but not a sparse solution.



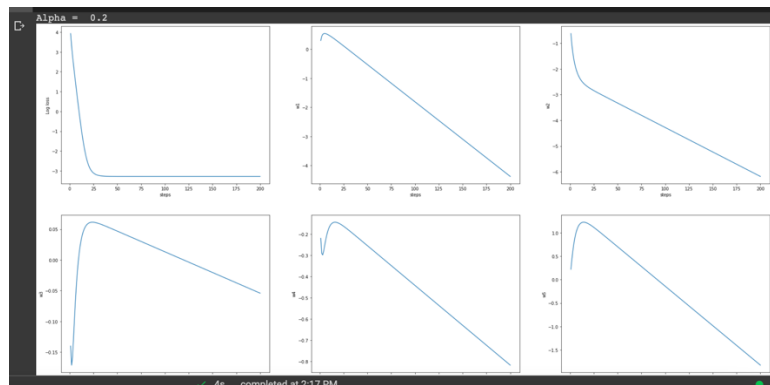
c. Problem 2.3:

From the plot below, we can see that W is converging to a value after several steps are executed. Yes, W is converging to an optimal solution and a sparse solution.

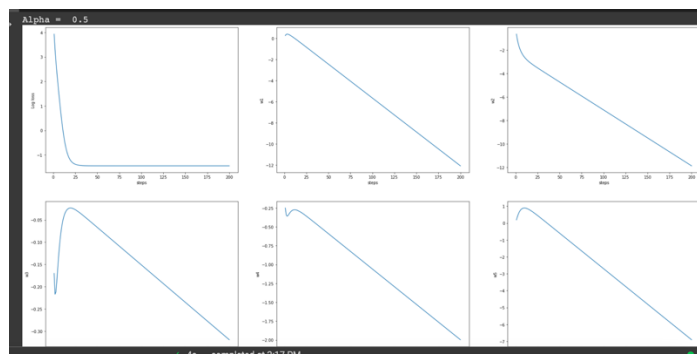


d. Problem 2.4:

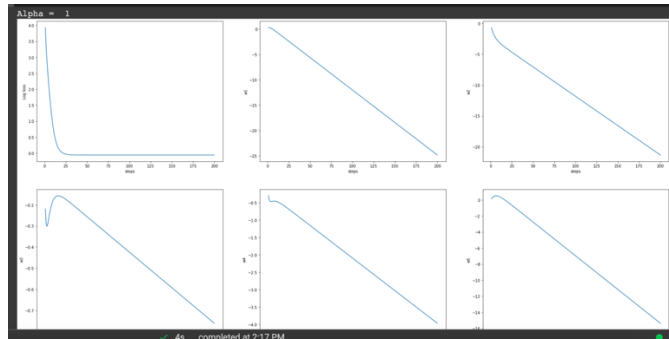
- i. $\alpha = 0.2$, from the plot we can see that loss converged to -3 after some steps and W hasn't reach a steady point after 200 steps



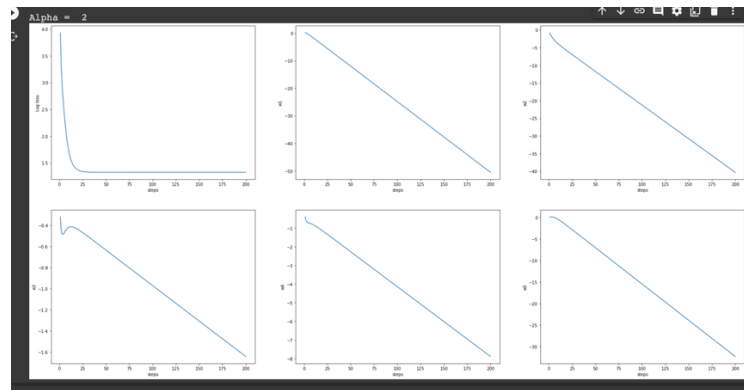
- ii. $\alpha = 0.5$, we can see that the loss is much lower the $\alpha=0.2$, however, the weight is still dropping and not yet reach a steady state.



- iii. $\alpha = 1$, loss converge at 0 and this might the best case, weight has been dropping drastically.




- iv. $\alpha = 2$, loss was greater than the previous alpha values. However, the weight turns into a steady drop, unlike the previous ones, which increases and decreases, when $\alpha = 2$, the weight dropped consistency.




e. Problem 2.5:


- i. Threshold = 0.004, we can see that comparing with 2.4, the loss is dropping consistency, whereas the weight is converging, though some of the data have a bump.




- ii. Threshold = 0.01, we can see that all the weight converges when reached step 200, however the loss is not converging.



- iii. Threshold = 0.02, all the weight had converged, and so does the loss, the performance was better than 2.4



- iv. Threshold = 0.04, we can see that the convergence of the weight reaches their steady point much faster than the previous ones. Loss, on the other hand, doesn't converge yet after 200 steps.



f. Problem 2.6:

3. Lab2: Pruning ResNet-20 model

a. Problem 3.1:

The accuracy of the floating-point pretrained model is 0.9151

```
# Load the best weight parameters
net.load_state_dict(torch.load("pretrained_model.pt"))
test(net)

Downloading https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz to ./data/cifar-10-python.tar.gz
100% 170498071/170498071 [00:05<00:00, 34187172.26it/s]
Extracting ./data/cifar-10-python.tar.gz to ./data
Test Loss=0.3231, Test accuracy=0.9151
```

b. Problem 3.2:

```
import numpy.ma as ma
def prune_by_percentage(layer, q=80.0):
    """
    Pruning the weight parameters by threshold.
    :param q: pruning percentile, 'q' percent of the least
    significant weight parameters will be pruned.
    """
    # Convert the weight of "layer" to numpy array
    new_layer = layer.weight.detach().cpu().numpy()

    # Compute the q-th percentile of the abs of the converted array
    temp_layer = np.absolute(new_layer)
    qth = np.percentile(temp_layer, q#, axis = 0, keepdims=False)

    # Generate a binary mask same shape as weight to decide which element to prune
    mask = np.where(temp_layer < qth, 0, new_layer)

    # Convert mask to torch tensor and put on GPU
    new_mask = torch.tensor(mask)

    # Multiply the weight by mask to perform pruning
    layer.weight.data = new_mask.to(device)
    pass
```

c. Problem 3.3:

From the result shown in the image, the best accuracy is 88.05% after 20 epochs of training.

```
Epoch: 18
[Step=50] Loss=0.2489 acc=0.9123 1635.6 examples/second
[Step=100] Loss=0.2460 acc=0.9136 2534.9 examples/second
[Step=150] Loss=0.2448 acc=0.9147 2553.7 examples/second
Test Loss=0.3675, Test acc=0.8793
Saving...

Epoch: 19
[Step=50] Loss=0.2500 acc=0.9128 1600.0 examples/second
[Step=100] Loss=0.2466 acc=0.9134 2570.4 examples/second
[Step=150] Loss=0.2412 acc=0.9148 2539.9 examples/second
Test Loss=0.3642, Test acc=0.8805
Saving...
```

d. Problem 3.4:

From the testing result, the best accuracy occurs at the first two epochs with accuracy = 91.55%, which is different with the previous result, which increases gradually.

```

Epoch: 0
qth = 8
[Step=50]      Loss=0.0494      acc=0.9840      1581.2 examples/second
[Step=100]     Loss=0.0496      acc=0.9839      2545.3 examples/second
[Step=150]     Loss=0.0496      acc=0.9840      2560.3 examples/second
Test Loss=0.3231, Test acc=0.9155

Epoch: 1
qth = 16
[Step=50]      Loss=0.0495      acc=0.9832      1604.1 examples/second
[Step=100]     Loss=0.0486      acc=0.9840      2551.5 examples/second
[Step=150]     Loss=0.0489      acc=0.9839      2495.6 examples/second
Test Loss=0.3268, Test acc=0.9140

Epoch: 2

```

e. Problem 3.5:

The result shows the same which the best accuracy occurs during the first couple epochs. The total sparsity was 79%.

```

Epoch: 0
q = 8
[Step=50]      Loss=0.0473      acc=0.9853      1556.6 examples/second
[Step=100]     Loss=0.0464      acc=0.9855      2590.2 examples/second
[Step=150]     Loss=0.0471      acc=0.9854      2541.4 examples/second
Test Loss=0.3254, Test acc=0.9145

Epoch: 1
q = 16
[Step=50]      Loss=0.0471      acc=0.9858      1610.5 examples/second
[Step=100]     Loss=0.0479      acc=0.9850      2594.2 examples/second
[Step=150]     Loss=0.0479      acc=0.9845      2566.5 examples/second
Test Loss=0.3285, Test acc=0.9146

```

```

Sparsity of head_conv.0.conv: 0.3101851851851852
Sparsity of body_op.0.conv1.0.conv: 0.6575520833333334
Sparsity of body_op.0.conv2.0.conv: 0.6380208333333334
Sparsity of body_op.1.conv1.0.conv: 0.6267361111111112
Sparsity of body_op.1.conv2.0.conv: 0.6493055555555556
Sparsity of body_op.2.conv1.0.conv: 0.6315104166666666
Sparsity of body_op.2.conv2.0.conv: 0.6684027777777778
Sparsity of body_op.3.conv1.0.conv: 0.6236979166666666
Sparsity of body_op.3.conv2.0.conv: 0.6883680555555556
Sparsity of body_op.4.conv1.0.conv: 0.7254774305555556
Sparsity of body_op.4.conv2.0.conv: 0.7829861111111112
Sparsity of body_op.5.conv1.0.conv: 0.7241753472222222
Sparsity of body_op.5.conv2.0.conv: 0.8132595486111112
Sparsity of body_op.6.conv1.0.conv: 0.7325303819444444
Sparsity of body_op.6.conv2.0.conv: 0.7645670572916666
Sparsity of body_op.7.conv1.0.conv: 0.7768825954861112
Sparsity of body_op.7.conv2.0.conv: 0.8259548611111112
Sparsity of body_op.8.conv1.0.conv: 0.8529730902777778
Sparsity of body_op.8.conv2.0.conv: 0.9767523871527778
Sparsity of final_fc.linear: 0.15625
Total sparsity of: 0.7999970186631685
Files already downloaded and verified
Test Loss=0.3473, Test accuracy=0.8834

```

4. Lab3: Fixed-point quantization and fine-tuning

a. Problem 4.1:

```

class STE(torch.autograd.Function):
    @staticmethod
    def forward(ctx, w, bit):
        if bit is None:
            wq = w
        elif bit==0:
            wq = w*0
        else:
            # For Lab 3 bounds only (optional), build a mask to record position of zero weights
            #weight_mask = ...

            # Lab3 (a), Your code here:
            # Compute alpha (scale) for dynamic scaling
            alpha = np.max(w.detach().cpu().numpy()) - np.min(w.detach().cpu().numpy())
            # Compute beta (bias) for dynamic scaling
            beta = np.min(w.detach().cpu().numpy())
            # Scale w with alpha and beta so that all elements in ws are between 0 and 1
            ws = (w - beta) / alpha

            step = 2 ** (bit)-1
            # Quantize ws with a linear quantizer to "bit" bits
            R = (1 / step) * torch.round(step * ws)
            # Scale the quantized weight R back with alpha and beta
            wq = alpha * R + beta

            # For Lab 3 bounds only (optional), restore zero elements in wq
            #wq = wq*weight_mask

        return wq

```

b. Problem 4.2:

From the result, we can see that when Nbits = 6, the accuracy is the best, and it decreases as we lower the Nbit value.

```

-----
current Nbits = 6
Files already downloaded and verified
Test Loss=0.3364, Test accuracy=0.9145
-----
current Nbits = 5
Files already downloaded and verified
Test Loss=0.3390, Test accuracy=0.9112
-----
current Nbits = 4
Files already downloaded and verified
Test Loss=0.3858, Test accuracy=0.8972
-----
current Nbits = 3
Files already downloaded and verified
Test Loss=0.9874, Test accuracy=0.7662
-----
current Nbits = 2
Files already downloaded and verified
Test Loss=9.5441, Test accuracy=0.0899

```

c. Problem 4.3:

As we can see from the following results, when the value of Nbit is greater, the accuracy is better, and so does the quantize results.

i. Nbits = 4

```

Epoch: 18
[Step=3550] Loss=0.0547 acc=0.9815 1179.0 examples/second
[Step=3600] Loss=0.0513 acc=0.9827 2337.2 examples/second
[Step=3650] Loss=0.0530 acc=0.9826 2321.9 examples/second
[Step=3700] Loss=0.0535 acc=0.9825 2494.7 examples/second
Test Loss=0.3357, Test acc=0.9127

Epoch: 19
[Step=3750] Loss=0.0542 acc=0.9811 1153.1 examples/second
[Step=3800] Loss=0.0538 acc=0.9815 2338.6 examples/second
[Step=3850] Loss=0.0554 acc=0.9812 2362.4 examples/second
[Step=3900] Loss=0.0547 acc=0.9813 2561.5 examples/second
Test Loss=0.3341, Test acc=0.9135
Saving...
Files already downloaded and verified
Test Loss=0.3341, Test accuracy=0.9135
-----
current Nbits = 3

```


ii. Nbits = 3

```
Epoch: 15
[Step=2950]    Loss=0.0955    acc=0.9672    1183.0 examples/second
[Step=3000]    Loss=0.0943    acc=0.9661    2367.6 examples/second
[Step=3050]    Loss=0.0911    acc=0.9670    2342.5 examples/second
[Step=3100]    Loss=0.0900    acc=0.9678    2403.1 examples/second
Test Loss=0.3545, Test acc=0.9049

Epoch: 16
[Step=3150]    Loss=0.0789    acc=0.9729    1210.4 examples/second
[Step=3200]    Loss=0.0888    acc=0.9691    2321.8 examples/second
[Step=3250]    Loss=0.0889    acc=0.9689    2327.1 examples/second
[Step=3300]    Loss=0.0884    acc=0.9687    2335.9 examples/second
Test Loss=0.3513, Test acc=0.9076
Saving...

Epoch: 17
```

iii. Nbits = 2

```
Test Loss=0.4702, Test acc=0.8506

Epoch: 14
[Step=2750]    Loss=0.3100    acc=0.8900    1191.9 examples/second
[Step=2800]    Loss=0.3002    acc=0.8940    2365.7 examples/second
[Step=2850]    Loss=0.3004    acc=0.8950    2324.6 examples/second
[Step=2900]    Loss=0.3026    acc=0.8945    2337.3 examples/second
Test Loss=0.4701, Test acc=0.8499

Epoch: 15
[Step=2950]    Loss=0.3017    acc=0.8945    1191.8 examples/second
[Step=3000]    Loss=0.2951    acc=0.8962    2365.2 examples/second
[Step=3050]    Loss=0.2993    acc=0.8939    2384.7 examples/second
[Step=3100]    Loss=0.2983    acc=0.8951    2336.7 examples/second
Test Loss=0.4652, Test acc=0.8565
Saving...

Epoch: 16
[Step=3150]    Loss=0.2935    acc=0.8954    1204.8 examples/second
[Step=3200]    Loss=0.2924    acc=0.8990    2376.3 examples/second
```

d. Problem 4.4:

e. Problem 4.5: